

Chapter 1

Design

1.1 System Architecture

In this section, I will lay out the overall logic, and an overview of the steps involved in running my program. By decomposing the program into individual abstracted stages, I can focus on the workings and functionality of each section individually, which makes documenting and coding each section easier. I have also included a flowchart to illustrate the logic of each screen of the program.

I will also create an abstracted GUI prototype in order to showcase the general functionality of the user experience, while acting as a reference for further stages of graphical development. It will consist of individually drawn screens for each stage of the program, as shown in the top-level overview. The elements and layout of each screen are also documented below.

The following is a top-level overview of the logic of the program:

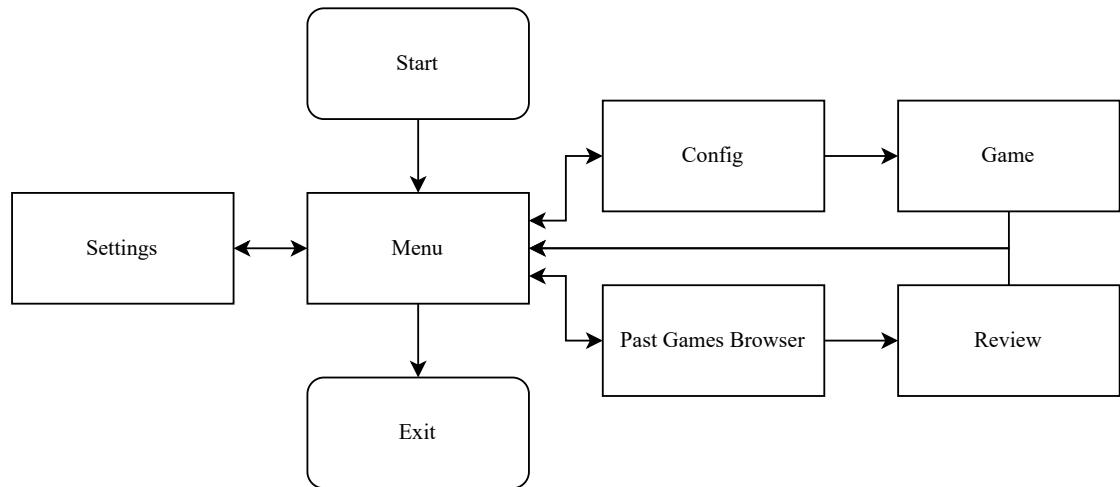


Figure 1.1: Flowchart for Program Overview

1.1.1 Main Menu



Figure 1.2: Main Menu screen prototype

The main menu will be the first screen to be displayed, providing access to different stages of the game. The GUI should be simple yet effective, containing clearly-labelled buttons for the user to navigate to different parts of the game.

1.1.2 Settings

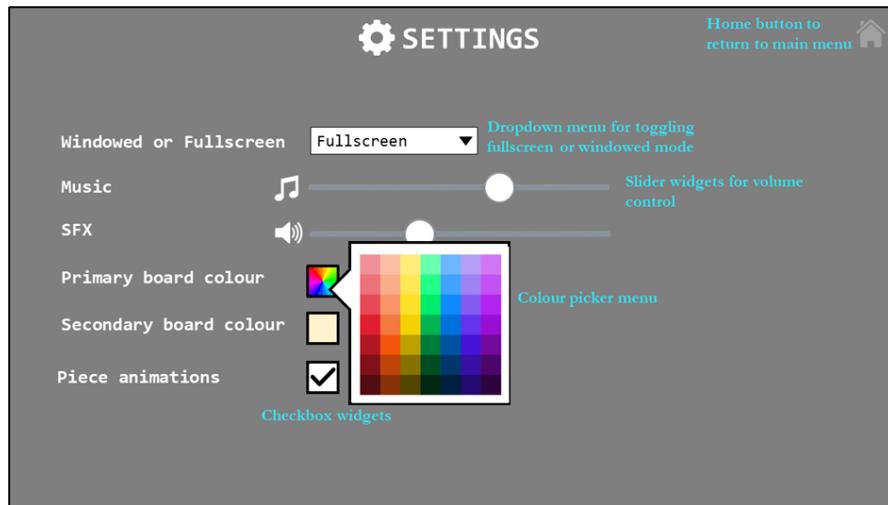


Figure 1.3: Settings screen prototype

The settings menu allows for the user to customise settings related to the program as a whole. The settings will be changed via GUI elements such as buttons and sliders, offering the ability

to customize display mode, volume, board colour etc. Changes to settings will be stored in an intermediate code class, then stored externally into a JSON file. Game settings will instead be changed in the Config screen.

The setting screen should provide a user-friendly interface for changing the program settings intuitively; I have therefore selected appropriate GUI widgets for each setting:

- Windowed or Fullscreen - Drop-down list for selecting between pre-defined options
- Music and SFX - Slider for selecting audio volume, a continuous value
- Board colour - Colour grid for the provision of multiple pre-selected colours
- Piece animation - Checkbox for toggling between on or off

Additionally, each screen is provided with a home button icon on the top right (except the main menu), as a shortcut to return to the main menu.

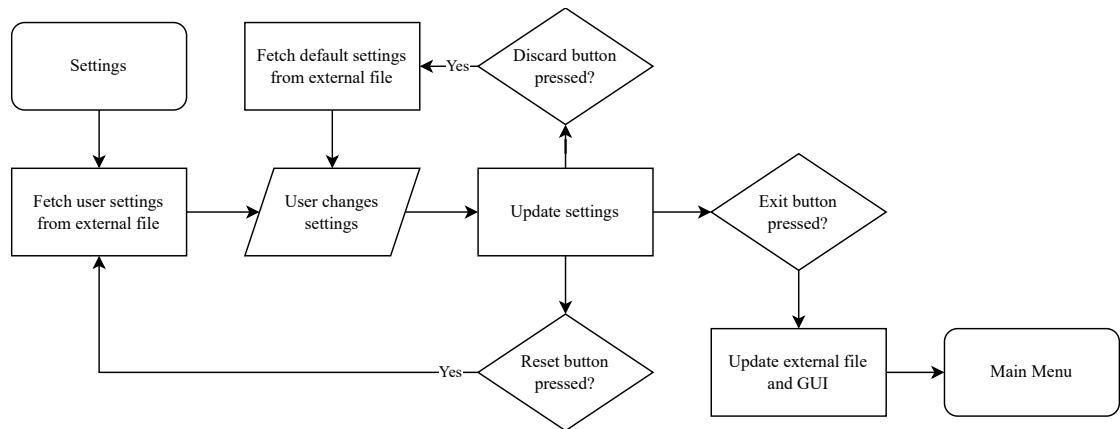


Figure 1.4: Flowchart for Settings

1.1.3 Past Games Browser



Figure 1.5: Browser screen prototype

The Past Games Browser menu displays a list of previously played games to be replayed. When selecting a game, the replay will render out the saved FEN string into a board position identical to the one played previously, except the user is limited to replaying back and forth between recorded moves. The menu also offers the functionality of sorting games in terms of time, game length etc.

For the GUI, previous games will be displayed on a strip, scrolled through by a horizontal slider. Information about the game will be displayed for each instance, along with the option to copy the FEN string to be stored locally or to be entered into the Review screen. When choosing a past game, a green border will appear to show the current selection, and double clicking enters the user into the full replay mode. While replaying the game, the GUI will appear identical to an actual game. However, the user will be limited to scrolling throughout the moves via the left and right arrow keys.

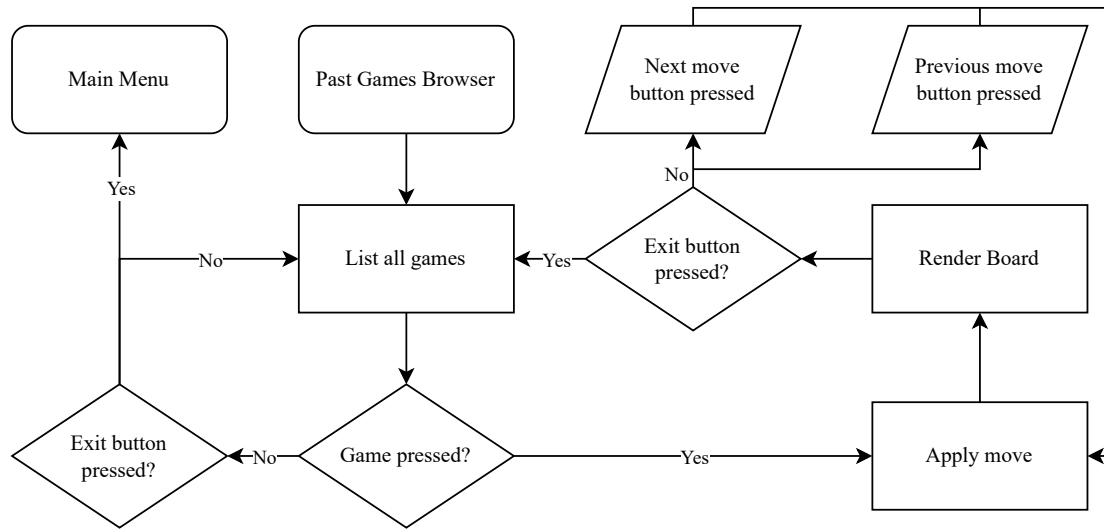


Figure 1.6: Flowchart for Browser

1.1.4 Config

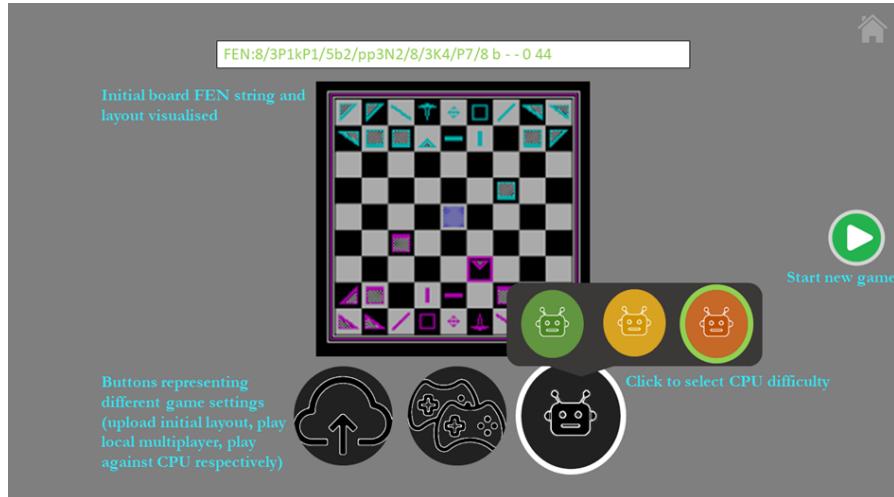


Figure 1.7: Config screen prototype

The config screen comes prior to the actual gameplay screen. Here, the player will be able to change game settings such as toggling the CPU player, time duration, playing as white or black etc.

The config menu is loaded with the default starting position. However, players may enter their own FEN string as an initial position, with the central board updating responsively to give a visual representation of the layout. Players are presented with the additional options to play against a friend, or against a CPU, which displays a drop-down list when pressed to select the CPU difficulty.

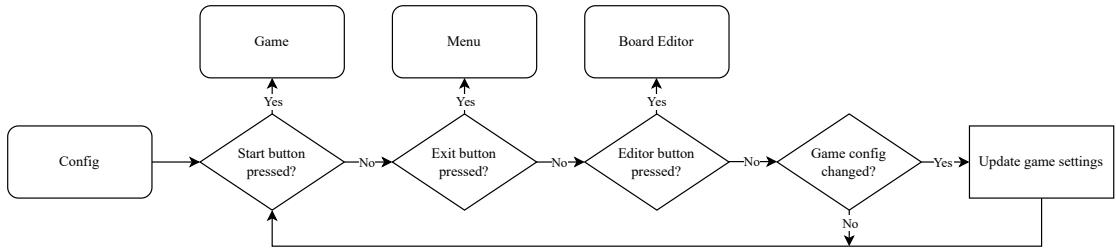


Figure 1.8: Flowchart for Config

1.1.5 Game



Figure 1.9: Game screen prototype

During the game, handling of the game logic, such as calculating player turn, calculating CPU moves or laser trajectory, will be computed by the program internally, rendering the updated GUI accordingly in a responsive manner to provide a seamless user experience.

In the game screen, the board is positioned centrally on the screen, surrounded by accompanying widgets displaying information on the current state of the game. The main elements include:

- Status text - displays information on the game state and prompts for each player move
- Rotation buttons - allows each player to rotate the selected piece by 90° for their move
- Timer - displays available time left for each player
- Draw and forfeit buttons - for the named functionalities, confirmed by pressing twice
- Piece display - displays material captured from the opponent for each player

Additionally, the current selected piece will be highlighted, and the available squares to move to will also contain a circular visual cue. Pieces will either be moved by clicking the

target square, or via a drag-and-drop mechanism, accompanied by responsive audio cues. These implementations aim to improve user-friendliness and intuitiveness of the program.

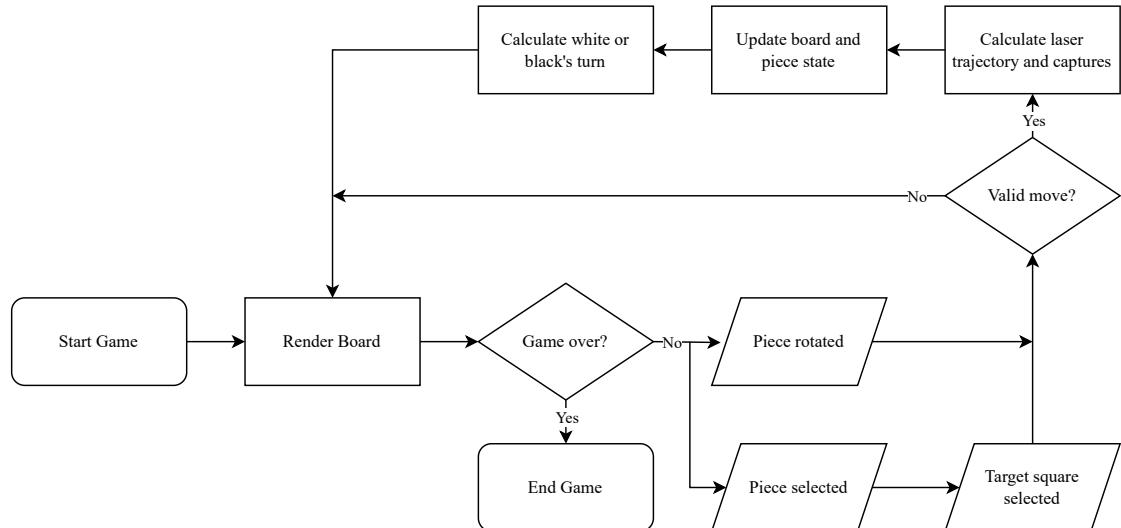


Figure 1.10: Flowchart for Game

1.1.6 Board Editor

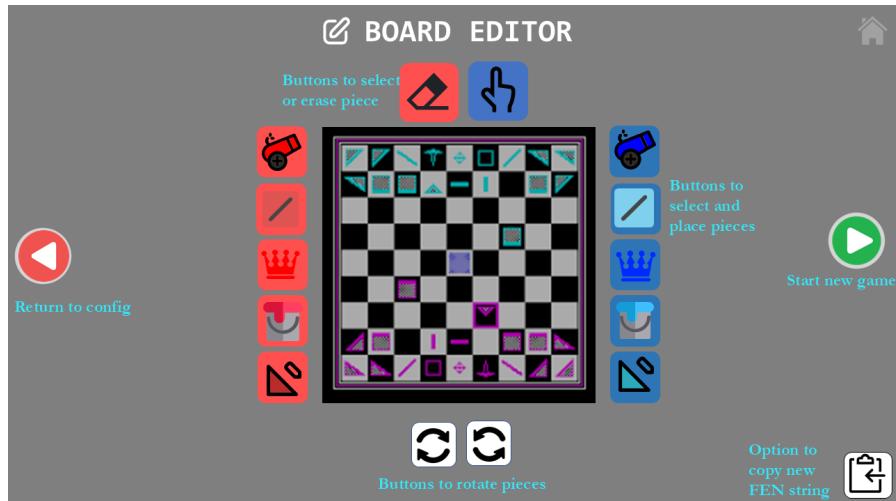


Figure 1.11: Editor screen prototype

The editor screen is used to configure the starting position of the board. Controls should include the ability to place all piece types of either colour, to erase pieces, and easy board manipulation shortcuts such as dragging pieces or emptying the board.

For the GUI, the buttons should clearly represent their functionality, through the use of icons and appropriate colouring (e.g. red for delete).

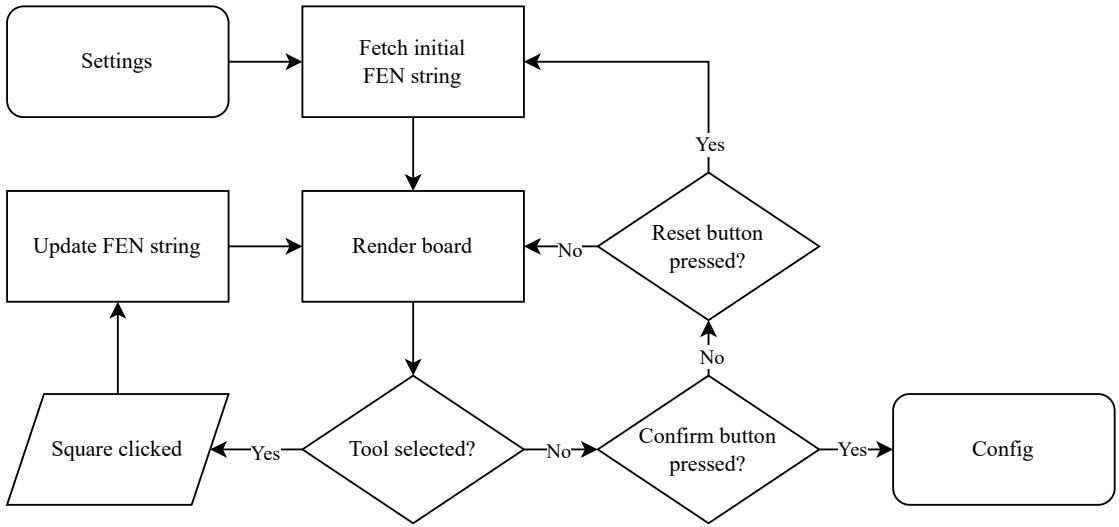


Figure 1.12: Flowchart for board editor

1.2 Algorithms and Techniques

1.2.1 Minimax

Minimax is a backtracking algorithm commonly used in zero-sum games used to determine the score according to an evaluation function, after a certain number of perfect moves. Minimax aims to minimize the maximum advantage possible for the opponent, thereby minimizing a player's possible loss in a worst-case scenario. It is implemented using a recursive depth-first search, alternating between minimizing and maximizing the player's advantage in each recursive call.

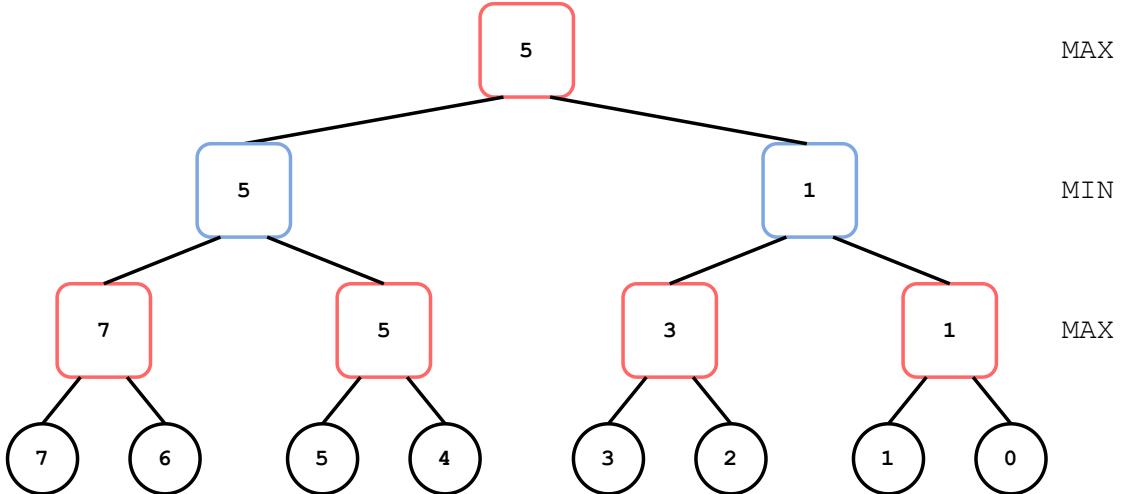


Figure 1.13: Example minimax tree

For the example minimax tree show in Figure 1.13, starting from the bottom leaf node

evaluations, the maximising player would choose the highest values (7, 5, 3, 1). From those values, the minimizing player would choose the lowest values (5, 1). The final value chosen by the maximum player would therefore be the highest of the two, 5.

Implementation in the form of pseudocode is shown below:

Algorithm 1 Minimax pseudocode

```

function MINIMAX(node, depth, maximisingPlayer)
    if depth = 0 OR node equals game over then
        return EVALUATE
    end if

    if maximisingPlayer then
        value  $\leftarrow -\infty$ 
        for child of node do
            value  $\leftarrow \text{MAX}(value, \text{MINIMAX}(child, depth - 1, false))$ 
        end for
        return value
    else
        value  $\leftarrow +\infty$ 
        for child of node do
            value  $\leftarrow \text{MIN}(value, \text{MINIMAX}(child, depth - 1, true))$ 
        end for
        return value
    end if
end function

```

1.2.2 Minimax improvements

Alpha-beta pruning

Alpha-beta pruning is a search algorithm that aims to decrease the number of nodes evaluated by the minimax algorithm. Alpha-beta pruning stops evaluating a move in the game tree when one refutation is found in its child nodes, proving the node to be worse than previously-examined alternatives. It does this without any potential of pruning away a better move. The algorithm maintains two values: alpha and beta. Alpha (α), the upper bound, is the highest value that the maximising player is guaranteed of; Beta (β), the lower bound, is the lowest value that the minimizing player is guaranteed of. If the condition $\alpha \geq \beta$ for a node being evaluated, the evaluation process halts and its remaining children nodes are ‘pruned’.

This is shown in the following maximising example:

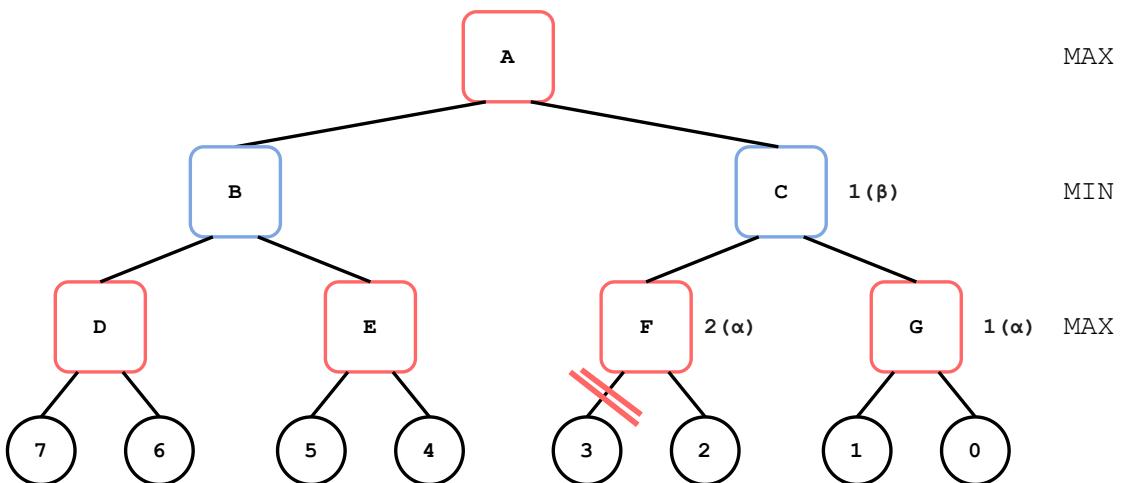


Figure 1.14: Example minimax tree with alpha-beta pruning

Since minimax is a depth-first search algorithm, nodes C and G and their α and β have already been searched. Next, at node F , the current α and β are $-\infty$ and 1 respectively, since the β is passed down from node C . Searching the first leaf node, the α subsequently becomes $\alpha = \max(-\infty, 2)$. This means that the maximising player at this depth is already guaranteed an evaluation of 2 or greater. Since we know that the minimising player at the depth above is guaranteed a value of 1, there is no point in continuing to search node F , a node that returns a value of 2 or greater. Hence at node F , where $\alpha \geq \beta$, the branches are pruned.

Alpha-beta pruning therefore prunes insignificant nodes by maintain an upper bound α and lower bound β . This is an essential optimization as a simple minimax tree increases exponentially in size with each depth ($O(b^d)$, with branching factor b and d ply depth), and alpha-beta reduces this and the associated computational time considerably.

The pseudocode implementation is shown below:

Algorithm 2 Minimax with alpha-beta pruning pseudocode

```
function MINIMAX(node, depth,  $\alpha$ ,  $\beta$ , maximisingPlayer)
    if  $depth = 0$  OR node equals game over then
        return EVALUATE
    end if

    if maximisingPlayer then
        value  $\leftarrow -\infty$ 
        for child of node do
            value  $\leftarrow \text{MAX}(value, \text{MINIMAX}(child, depth - 1, \alpha, \beta, false))$ 
            if  $value > \beta$  then break
        end if
         $\alpha \leftarrow \text{MAX}(\alpha, value)$ 
    end for
    return value
else
    value  $\leftarrow +\infty$ 
    for child of node do
        value  $\leftarrow \text{MIN}(value, \text{MINIMAX}(child, depth - 1, \alpha, \beta, true))$ 
        if  $value < \alpha$  then break
    end if
     $\beta \leftarrow \text{MIN}(\beta, value)$ 
end for
return value
end if
end function
```

Transposition Tables & Zobrist Hashing

Transition tables, a memoisation technique, again greatly reduces the number of moves searched. During a brute-force minimax search with a depth greater than 1, the same positions may be searched multiple times, as the same position can be reached from different sequences of moves. A transposition table caches these same positions (transpositions), along with its associated evaluations, meaning commonly reached positions are not unnecessarily re-searched.

Flags and depth are also stored alongside the evaluation. Depth is required as if the current search comes across a cached position with an evaluation calculated at a lower depth than the current search, the evaluation may be inaccurate. Flags are required for dealing with the uncertainty involved with alpha-beta pruning, and can be any of the following three.

Exact flag is used when a node is fully searched without pruning, and the stored and fetched evaluation is accurate.

Lower flag is stored when a node receives an evaluation greater than the β , and is subsequently pruned, meaning that the true evaluation could be higher than the value stored. We are thus storing the α and not an exact value. Thus, when we fetch the cached value, we have to recheck if this value is greater than β . If so, we return the value and this branch is pruned (fail high); If not, nothing is returned, and the exact evaluation is calculated.

Upper flag is stored when a node receives an evaluation smaller than the α , and is subsequently pruned, meaning that the true evaluation could be lower than the value stored. Similarly, when we fetch the cached value, we have to recheck if this value is lower than α . Again, the current branch is pruned if so (fail low), and an exact evaluation is calculated if not.

The pseudocode implementation for transposition tables is shown below:

Algorithm 3 Minimax with transposition table pseudocode

```

function MINIMAX(node, depth,  $\alpha$ ,  $\beta$ , maximisingPlayer)
    hash_key  $\leftarrow$  HASH(node)
    entry  $\leftarrow$  GETENTRY(hash_key)

    if entry.hash_key = hash_key AND entry.hash_key  $\geq$  depth then
        if entry.hash_key = EXACT then
            return entry.value
        else if entry.hash_key = LOWER then
             $\alpha \leftarrow \text{MAX}(\alpha, \text{entry.value})$ 
        else if entry.hash_key = UPPER then
             $\beta \leftarrow \text{MIN}(\beta, \text{entry.value})$ 
        end if
        if  $\alpha \geq \beta$  then
            return entry.value
        end if
    end if

    ...normal minimax...

    entry.value  $\leftarrow$  value
    entry.depth  $\leftarrow$  depth
    if value  $\leq \alpha$  then
        entry.flag  $\leftarrow$  UPPER
    else if value  $\geq \beta$  then
        entry.flag  $\leftarrow$  LOWER
    else
        entry.flag  $\leftarrow$  EXACT
    end if

    return value
end function

```

The current board position will be used as the index for a transposition table entry. To convert our board state and bitboards into a valid index, Zobrist hashing may be used. For every square on the chessboard, a random integer is assigned to every piece type (12 in our case, 6 piece type, times 2 for both colours). To initialise a hash, the random integer associated with the piece on a specific square undergoes a XOR operation with the existing hash. The hash is incrementally update with XOR operations every move, instead of being recalculated from scratch improving computational efficiency. Using XOR operations also allows moves to be reversed, proving useful for the functionality to scroll through previous moves. A Zobrist hash is also a better candidate than FEN strings in checking for threefold-repetition, as they are less

intensive to calculate for every move.

The pseudocode implementation for Zobrist hashing is shown below:

Algorithm 4 Zobrist hashing pseudocode

RANDOMINTS represents a pre-initialised array of random integers for each piece type for each square

```
function HASH _ BOARD(board)
    hash ← 0
    for each square on board do
        if square is not empty then
            hash ⊕ RANDOMINTS[square][piece on square]
        end if
    end for
    return hash
end function

function UPDATEHASH(hash, move)
    hash ⊕ RANDOMINTS[source square][piece]
    hash ⊕ RANDOMINTS[destination square][piece]
    if red to move then
        hash ⊕ hash for red to move ▷ Hash needed for move colour, as two identical positions
        are different if the colour to move is different
    end if
    return hash
end function
```

1.2.3 Board Representation

FEN string

Forsyth-Edwards Notation (FEN) notation provides all information on a particular position in a chess game. I intend to implement methods parsing and generating FEN strings in my program, in order to load desired starting positions and save games for later play. Deviating from the classic 6-part format, a custom FEN string format will be required for our laser chess game, accommodating its different rules from normal chess.

Our custom format implementation is show by the example below:

sc3ncfancpb2/2pc7/3Pd7/pa1Pc1rbra1pb1Pd/pb1Pd1RaRb1pa1Pc/6pb3/7Pa2/2PdNaFaNa3Sa
r

Our FEN string format contains two parts, denoted by the space between them:

- Part 1: Describes the location of each piece. The construction of this part is defined by the following rules:
 - The board is read from top-left to bottom-right, row by row
 - A number represents the number of empty squares before the next piece
 - A capital letter represents a blue piece, and a lowercase letter represents a red piece

- The letters F , R , P , N , S stand for the pieces Pharaoh, Scarab, Pyramid, Anubis and Sphinx respectively
- Each piece letter is followed by the lowercase letters a , b , c or d , representing a 0° , 90° , 180° and 270° degree rotation respectively
- Part 2: States the active colour, b means blue to move, r means red to move

Having inputted the desired FEN string board configuration in the config menu, the bitboards for each piece will be initialised with the following functions:

Algorithm 5 FEN string pseudocode

```

function PARSE_FEN_STRING(fen_string, board)
    part_1, part_2  $\leftarrow$  SPLIT(fen_string)
    rank  $\leftarrow$  8
    file  $\leftarrow$  0

    for character in part_1 do
        square  $\leftarrow$  rank  $\times$  8 + file
        if character is alphabetic then
            if character is lower then
                board.bitboards[red][character]  $\mid\mid$  1 << character
            else
                board.bitboards[blue][character]  $\mid\mid$  1 << character
            end if
        else if character is numeric then
            file  $\leftarrow$  file + character
        else if character is / then
            rank  $\leftarrow$  rank - 1
            file  $\leftarrow$  file + 1
        else
            file  $\leftarrow$  file + 1
        end if

        if part_2 is b then
            board.active_colour  $\leftarrow$  b
        else
            board.active_colour  $\leftarrow$  r
        end if
    end for
end function

```

The function first processes every piece and corresponding square in the FEN string, modifying each piece bitboard using a bitwise OR operator, with a 1 shifted over to the correctly occupied square using a Left-Shift operator. For the second part, the active colour property of the board class is initialised to the correct player.

Bitboards

Bitboards are an array of bits representing a position or state of a board game. Multiple bitboards are used with each representing a different property of the game (e.g. scarab position and

scarab rotation), and can be masked together or transformed to answer queries about positions. Bitboards offer an efficient board representation, its performance primarily arising from the speed of parallel bitwise operations used to transform bitboards. To map each board square to a bit in each number, we will assign each square from left to right, with the least significant bit (LSB) assigned to the bottom-left square (A1), and the most significant bit (MSB) to the top-right square (J8).

8	70	71	72	73	74	75	76	77	78	79
7	60	61	62	63	64	65	66	67	68	69
6	50	51	52	53	54	55	56	57	58	59
5	40	41	42	43	44	45	46	47	48	49
4	30	31	32	33	34	35	36	37	38	39
3	20	21	22	23	24	25	26	27	28	29
2	10	11	12	13	14	15	16	17	18	19
1	0	1	2	3	4	5	6	7	8	9
	a	b	c	d	e	f	g	h	j	k

Figure 1.15: Square to bit position mapping

Firstly, we need to initialise each bitboard and place 1s in the correct squares occupied by pieces. This is achieved whilst parsing the FEN-string, as shown in Algorithm 5. Secondly, we should implement an approach to calculate possible moves using our computed bitboards. We can begin by producing a bitboard containing the locations of all pieces, achieved through combining every piece bitboard with bitwise OR operations:

```
all_pieces_bitboard = white_pharaoh_bitboard | black_pharaoh_bitboard |
                     white_scarab_bitboard ...
```

Now, we can utilize this aggregated bitboard to calculate possible positional moves for each piece. For each piece, we can shift the entire bitboard to an adjacent target square (since every piece can only move one adjacent square per turn), and perform a bitwise AND operator with the bitboard containing all pieces, to determine if the target square is already occupied by an existing piece. For example, if we want to compute if the square to the left of our selected piece is available to move to, we will first shift every bit right (as the lowest square index is the LSB on the right, see diagram above), as demonstrated in the following 5x5 example:

	1	0		

Figure 1.16: `shifted_bitboard = piece_bitboard >> 1`

Where green represents the target square shifted into, and orange where the piece used to be. We can then perform a bitwise AND operation with the complement of the all pieces bitboard, where a square with a result of 1 represents an available target square to move to.

$$\text{available_squares_right} = (\text{piece_bitboard} >> 1) \& \sim \text{all_pieces_bitboard}$$

However, if the piece is on the leftmost A file, and is shifted to the right, it will be teleported onto the J file on the rank below, which is not a valid move. To prevent these erroneous moves for pieces on the edge of the board, we can utilise an A file mask to mask away any valid moves, as demonstrated below:

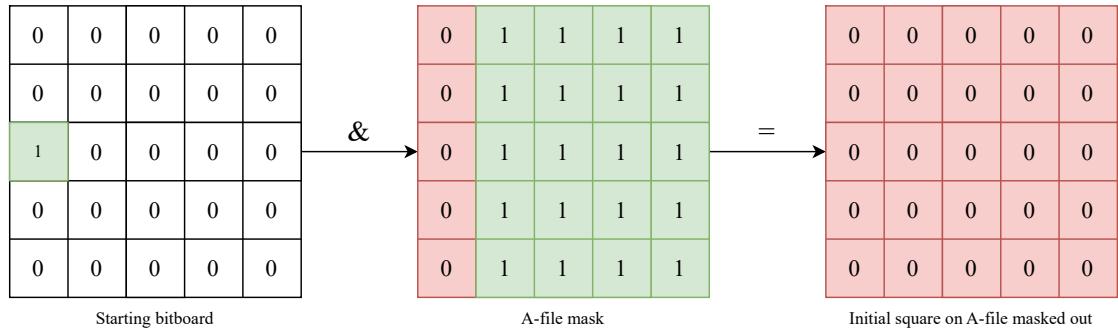


Figure 1.17: A-file mask example

This approach uses the logic that a piece on the A file can never move to a square on the left. Therefore, when calculating if a piece can move to a square on the left, we apply a bitwise AND operator with a mask where every square on the A file is 0; If a piece was on the A file, it will become 0, leaving no possible target squares to move to. The same approach can be mirrored for the far-right J file.

In theory, we do not need to implement the same solution for clipping in regards to ranks, as they are removed automatically by overflow or underflow when shifting bits too far. Our final function to calculate valid moves combines all the logic above: Shifting the selected piece in all

9 adjacent directions by their corresponding bits, masking away pieces trying to move into the edge of the board, combining them with a bitwise OR operator, and finally masking it with the all pieces bitboard to detect which squares are not currently occupied:

Algorithm 6 Finding valid moves pseudocode

```

function FIND_VALID_MOVES(selected_square)
    masked_a_square  $\leftarrow$  selected_square & A_FILE_MASK
    masked_j_square  $\leftarrow$  selected_square & J_FILE_MASK

    top_left  $\leftarrow$  masked_a_square << 9
    top_left  $\leftarrow$  masked_a_square << 9
    top_middle  $\leftarrow$  selected_square << 10
    top_right  $\leftarrow$  masked_<< 11
    middle_right  $\leftarrow$  masked_<< 1
    bottom_right  $\leftarrow$  masked_>> 9
    bottom_middle  $\leftarrow$  selected_square >> 10
    bottom_left  $\leftarrow$  masked_a_square >> 11
    middle_left  $\leftarrow$  masked_a_square >> 1

    possible_moves = top_left | top_middle | top_right | middle_right | bottom_right |
    bottom_middle | bottom_left | middle_left
    valid_moves = possible_moves & ~ALL_PIECES_BITBOARD

    return valid_moves
end function

```

1.2.4 Evaluation Function

The evaluation function is a heuristic algorithm to determine the relative value of a position. It outputs a real number corresponding to the advantage given to a player if reaching the analysed position, usually at a leaf node in the minimax tree. The evaluation function therefore provides the values on which minimax works on to compute an optimal move.

In the majority of evaluation functions, the most significant factor determining the evaluation is the material balance, or summation of values of the pieces. The hand-crafted evaluation function is then optimised by tuning various other positional weighted terms, such as board control and king safety.

Material Value

Since laser chess is not widely documented, I have assigned relative strength values to each piece according to my experience playing the game:

- Pharaoh - ∞
- Scarab - 200
- Anubis - 110
- Pyramid - 100

To find the number of pieces, we can iterate through the piece bitboard with the following popcount function:

Algorithm 7 Popcount pseudocode

```

function POPCOUNT(bitboard)
    count ← 0
    while bitboard do
        count ← count + 1
        bitboard ← bitboard&(bitboard - 1)
    end while
    return count
end function

```

Algorithm 7 continually resets the left-most 1 bit, incrementing a counter for each loop. Once the number of pieces has been established, we multiply this number by the piece value. Repeating this for every piece type, we can thus obtain a value for the total piece value on the board.

Piece-Square Tables

A piece in normal chess can differ in strength based on what square it is occupying. For example, a knight near the center of the board, controlling many squares, is stronger than a knight on the rim. Similarly, we can implement positional value for Laser Chess through Piece-Square Tables. PSQTs are one-dimensional arrays, with each item representing a value for a piece type on that specific square, encoding both material value and positional simultaneously. Each array will consist of 80 base values representing the piece's material value, with a bonus or penalty added on top for the location of the piece on each square. For example, the following PSQT is for the pharaoh piece type on an example 5x5 board:

0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Piece index

-10	-10	-10	-10	-10
-10	-10	-10	-10	-10
-5	-5	-5	-5	-5
0	0	0	0	0
5	5	5	5	5

Used to reference positional value in PSQT

Figure 1.18: PSQT showing the bonus position value gained for the square occupied by a pharaoh

For asymmetrical PSQTs, we would ideally like to label the board identically from both player's point of views. Although the PSQTs are displayed from the blue player's perspective (blue pharaoh at the bottom of the board), it uses indexes from the red player's perspective, as arrays and lists are defined with index 0 being at the topleft of the board. We would like to flip

the PSQTs to be reused with blue indexes, so that a generic algorithm can be used to sum up and calculate the total positional values for both players.

To utilise a PSQT for blue pieces, a special ‘FLIP’ table can be implemented:

8	70	71	72	73	74	75	76	77	78	79
7	60	61	62	63	64	65	66	67	68	69
6	50	51	52	53	54	55	56	57	58	59
5	40	41	42	43	44	45	46	47	48	49
4	30	31	32	33	34	35	36	37	38	39
3	20	21	22	23	24	25	26	27	28	29
2	10	11	12	13	14	15	16	17	18	19
1	0	1	2	3	4	5	6	7	8	9
	a	b	c	d	e	f	g	h	j	k

Figure 1.19: FLIP table used to map a blue piece index to the red player’s perspective

The FLIP table is just an array of indexes, mapping every blue player’s index onto the corresponding red index. The following expression utilises the FLIP table to retrieve a blue player’s value from the red player’s PSQT:

$$\text{blue_psqt_value} = \text{PHAROAH_PSQT}[\text{FLIP}[\text{square}]]$$

The following function retrieves an array of bitboards representing piece positions from the board class, then sums up all the values of these pieces for both players, referencing the corresponding PSQT:

Algorithm 8 Calculating positional value pseudocode

```

function CALCULATE_POSITIONAL_VALUE(bitboards, colour)
    positional_score ← 0
    for all pieces do
        for square in bitboards[piece] do
            if square = 1 then
                if colour is blue then
                    positional_score ← positional_score + PSQT[piece][square]
                else
                    positional_score ← positional_score + PSQT[piece][FLIP[square]]
                end if
            end if
        end for
    end for
    return positional_score
end function

```

Using valid squares

Using Algorithm 6 for finding valid moves, we can implement two more improvements for our evaluation function: Mobility and King Safety.

Mobility is the number of legal moves a player has for a given position. This is advantageous in most cases, with a positive correlation between mobility and the strength of a position. To implement this, we simply loop over all pieces of the active colour, and sum up the number of valid moves obtained from the previous algorithm.

King safety (Pharaoh safety) describes the level of protection of the pharaoh, being the piece that determines a win or loss. In normal chess, this would be achieved usually by castling, or protection via position or with other pieces. Similarly, since the only way to lose in Laser Chess is via a laser, having pieces surrounding the pharaoh, either to reflect the laser or to be sacrificed, is a sensible tactic and improves king safety. Thus, a value for king safety can be achieved by finding the number of valid moves a pharaoh can make, and subtracting them from the maximum possible of moves (8) to find the number of surrounding pieces.

1.2.5 Shadow Mapping

Following the client's requirement for engaging visuals, I have decided to implement shadow mapping for my program, especially as lasers are the main focus of the game. Shadow mapping is a technique used to create graphical hard shadows, with the use of a depth buffer map. I have chosen to implement shadow mapping, instead of alternative lighting techniques such as ray casting and ray marching, as its efficiency is more suitable for real-time usage, and results are visually decent enough for my purposes.

For typical 3D shadow mapping, the standard approach is as follows:

1. Render the scene from the light's point of view
2. Extract a depth buffer texture from the render
3. Compare the distance of a pixel from the light to the value stored in the depth texture
4. If greater, there must be an obstacle in the way reducing the depth map value, therefore that pixel must be in shadow

To implement shadow casting for my 2D game, I have modified some steps and arrived on the final following workflow:

1. Render the scene with only occluding objects shown
2. Crop texture to align the center to the light position
3. To create a 1D depth map, transform Cartesian to polar coordinates, and increase the distance from the origin until a collision with an occluding object
4. Using polar coordinates for the real texture, compare the z-depth to the corresponding value from the depth map
5. Additively blend the light colour if z-depth is less than the depth map value

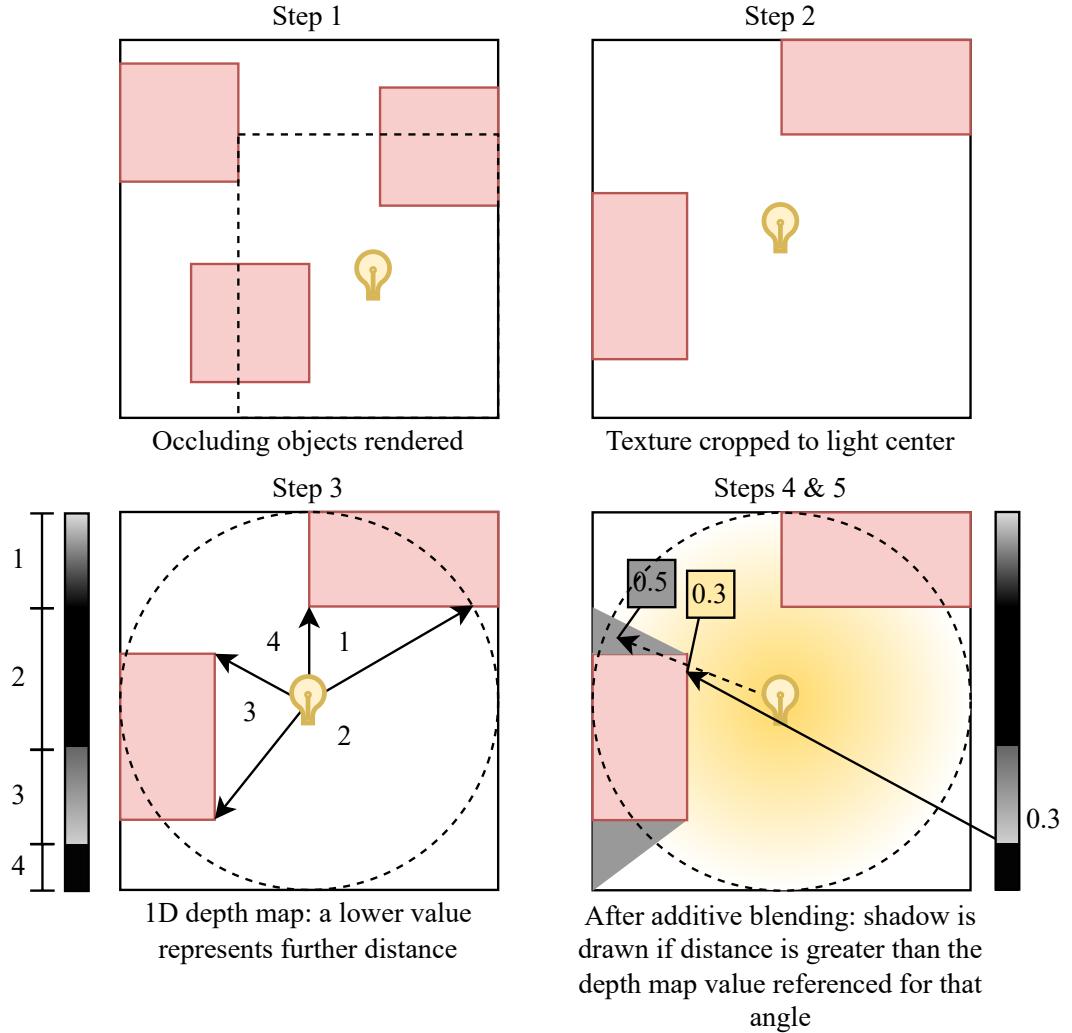


Figure 1.20: Workflow for 2D shadow mapping

Our method requires a coordinate transformation from Cartesian to polar, and vice versa. Polar to Cartesian transformation can be achieved with trigonometry, forming a right-angled triangle in the center and using the following two equations:

$$x = r \cos(\theta)$$

$$y = r \sin(\theta)$$

Cartesian to polar can also similarly be achieved with the right-angled triangle, finding the radius with the Pythagorean theorem, and the angle with arctan. However, since the range of the arctan function is only a half-circle ($\frac{\pi}{2} < \theta < \frac{3\pi}{2}$), we will have to use the atan2 function, which accounts for the negative quadrants, or the following:

$$\theta = 2 \arctan \left(\frac{r - x}{y} \right)$$

There are several disadvantages to shadow mapping. The relevant ones for us are Aliasing and Shadow Acne:

Aliasing occurs when the texture size for the depth map is smaller than the light map, causing shadows to be scaled up and rendered with jagged edges.

Shadow Acne occurs when the depth from the depth map is so close to the light map value, that precision errors cause unnecessary shadows to be rendered.

These problems can be mitigated by increasing the size of the shadow map size. However, due to memory and hardware constraints, I will have to find a compromised resolution to balance both artifacting and acuity.

Soft Shadows

The approach above is used only for calculating hard shadows. However, in real-life scenarios, lights are not modelled as a single particle, but instead emitted from a wide light source. This creates an umbra and penumbra, resulting in soft shadows.

To emulate this in our game, we could calculate penumbra values with various methods, however, due to hardware constraints and simplicity again, I have chosen to use the following simpler method:

1. Sample the depth map multiple times, from various differing angles
2. Sum the results using a normal distribution
3. Blur the final result proportional to the length from the center

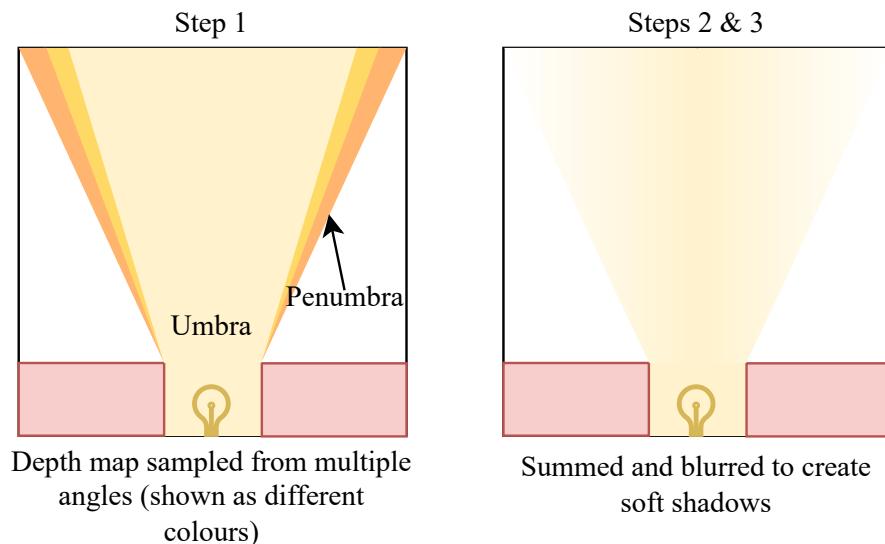


Figure 1.21: Workflow for 2D soft shadows

This method progressively blurs the shadow as the distance from the main shadow (umbra) increases, which results in a convincing estimation while being less computationally intensive.

1.2.6 Multithreading

In order to fulfill Objective ?? of a responsive GUI, I will have to employ multi-threading. Since python runs on a single thread natively, code is executed serially, meaning that a time consuming function such as minimax will prevent the running of another GUI-drawing function until it is finished, hence freezing the program. To overcome this, multi-threading can execute both functions in parallel on different threads, meaning the GUI-drawing thread can run while minimax is being computed, and stay responsive. To pass data between threads, since memory is shared between threads, arrays and queues can be used to store results from threads. The following flowchart shows my chosen approach to keep the GUI responsive while minimax is being computed:

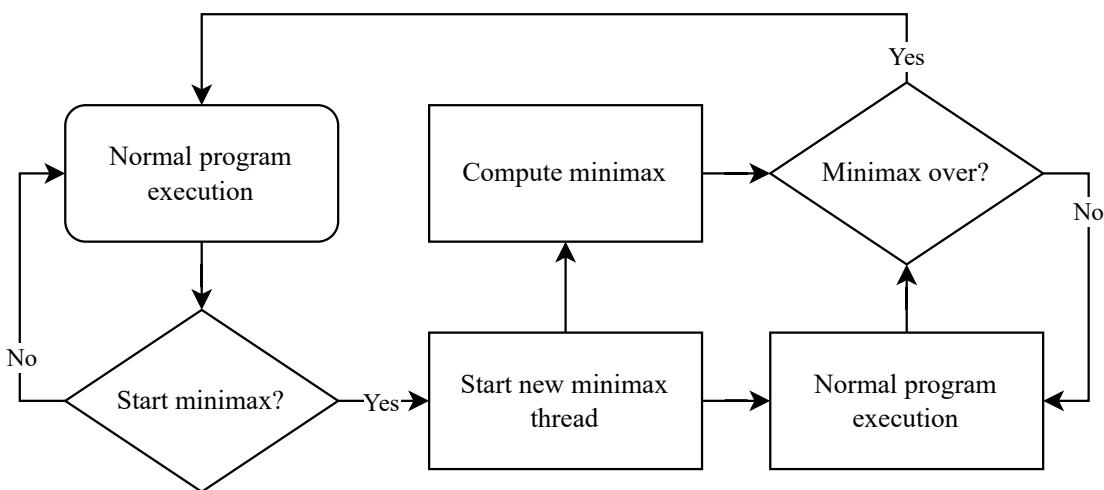


Figure 1.22: Multi-threading for minimax

1.3 Data Structures

1.3.1 Database

To achieve Objective ?? and stores previous games, I have opted to use a relational database. Choosing between different relational database, I have decided to use SQLite, since it does not require additional server softwars, has good performance with low memory requirements, and adequate for my use cases, with others such as Postgres being overkill.

DDL

Only a single entity will be required for my program, a table to store games. The table schema will be defined as follows:

Table: games

Field	Key	Data Type	Validation
game_id	Primary	INT	NOT NULL
winner		INT	

cpu_depth	INT	
number_of_moves	INT	NOT NULL
cpu_enabled	BOOL	NOT NULL
moves	TEXT	NOT NULL
initial_board_configuration	TEXT	NOT NULL
time	FLOAT	
created_dt	TIMESTAMP	NOT NULL

Table 1.1: Data table scheme for *games* table

All fields are either generated or retrieved from the board class, with the exception of the moves attribute, which will need to be encoded into a suitable data type such as a string. All attributes are also independent of each other¹, and so the the table therefore adheres to the third normal form.

To create the entity, a `CREATE` statement like the following can be used:

```

1 CREATE TABLE games(
2     id INTEGER PRIMARY KEY,
3     winner INTEGER,
4     cpu_depth INTEGER,
5     time real NOT NULL,
6     moves TEXT NOT NULL,
7     cpu_enabled INTEGER NOT NULL,
8     created_dt TIMESTAMP NOT NULL,
9     number_of_moves INTEGER NOT NULL,
10    initial_fen_string TEXT NOT NULL,
11 )

```

Removing an entity can also be done in a similar fashion:

```

1 DROP TABLE games

```

Migrations are a version control system to track incremental changes to the schema of a database. Since there is no popular SQL Python-binding libraries that support migrations, I will just be using a manual solution of creating python files that represent a change in my schema, defining functions that make use of SQL `ALTER` statements. This allows me to keep track of any changes, and rollback to a previous schema.

DML

To insert a new game entry into the table, an `INSERT` statement can be used with the provided array, where the appropiate arguments are binded to the correct attribute via `?` placeholders when run.

```

1 INSERT INTO games (
2     cpu_enabled,
3     cpu_depth,
4     winner,
5     time,
6     number_of_moves,
7     moves,

```

¹There is a case to be made for `moves` and `number_of_moves`, however I have included `number_of_moves` to save the computational effort of parsing the moves for every game just to display it on the browser preview section.

```

8     initial_fen_string ,
9     created_dt
10    )
11  VALUES  (?, ?, ?, ?, ?, ?, ?, ?, ?)

```

Moreover, we will need to fetch the number of total game entries in the table to be displayed to the user. To do this, the aggregate function `COUNT` can be used, which is supported by all SQL databases.

```
1  SELECT COUNT(*) FROM games
```

Pagination

When there are a large number of entries in the table, it would be appropriate to display all the games to the user in a paginated form, where they can scroll between different pages and groups of games. There are multiple methods to paginate data, such as using `LIMIT` and `OFFSET` clauses, or cursor-based pagination, but I have opted to use the `ROW_NUMBER()` function.

`ROW_NUMBER()` is a window function that assigns a sequential integer to a query's result set. If I were to query the entire table, each row would be assigned an integer that could be used to check if the row is in the bounds for the current page, and therefore be displayed. Moreover, the use of an `ORDER BY` clause enables sorting of the output rows, allowing the user to choose what order the games are presented in based on an attribute such as number of moves. A `PARTITION BY` clause will also be used to group the results base on an attribute such as winner prior to sorting, if the user wants to search for games based on multiple criteria with greater ease.

The start row and end row will be passed as parameters to the placeholders in the SQL statement, calculated by multiplying the page number by the number of games per page.

```

1  SELECT * FROM
2  (SELECT ROW_NUMBER() OVER (
3      PARTITION BY attribute1
4      ORDER BY attribute2 ASC
5  ) AS row_num, * FROM games)
6  WHERE row_num >= ? AND row_num <= ?

```

Security

Security measures such as database file permissions and encryption are common for a SQL database. However, since SQLite is a serverless database, and my program runs without any need for an internet connection, the risk of vulnerabilities is greatly reduced. Additionally, the game data stored on my database is frankly inconsequential, so going to great lengths to protect it wouldn't be to best use of my time. Nevertheless, my SQL Python-binding does support the user of placeholders for parameters, thereby addressing the risk of SQL injection attacks.

1.3.2 Linked Lists

Another data structure I intend to implement is linked lists. This will be integrated into widgets such as the carousel or multiple icon button widget, since these will contain a variable number of items, and where $O(1)$ random access is not a priority. Since moving back and forth between nodes is a must for a carousel widget, the linked list will be doubly-linked, with each node containing to its previous and next node. The list will also need to loop, with the next pointer of the last node pointing back to the first node, making it a circular linked list.

The following pseudocode outlines the basic functionality of the linked list:

Algorithm 9 Circular doubly linked list pseudocode

```
function INSERT_AT_FRONT(node)
    if head is none then
        head ← node
        node.next ← node.previous ← head
    else
        node.next ← head
        node.previous ← head.previous
        head.previous.next ← node
        head.previous ← node

        head ← node
    end if
end function
```

Require: $\text{LEN}(list) > 0$

```
function DATA_IN_LIST(data)
    current_node ← head.next
    while current_node ≠ head do
        if current_node.data = data then
            return True
        end if
        current_node ← current_node.next
    end while
    return False
end function
```

Require: Data in list

```
function REMOVE(data)
    current_node ← head
    while current_node.data ≠ data do
        current_node ← current_node.next
    end while

    current_node.previous.next ← current_node.next
    current_node.next.previous ← current_node.previous

    delete current_node
end function
```

1.3.3 Stack

Being a data structure with LIFO ordering, a stack is used for handling moves in the review screen. Starting with full stack of moves, every move undone pops an element off the stack to be processed. This move is then pushed onto a second stack. Therefore, cycling between moves requires pushing and popping between the two stacks, as shown in Figure 1.23. The same functionality can be achieved using a queue, but I have chosen to use two stacks as it is simpler

to implement, as being able to quickly check the number of items in each will come in handy.

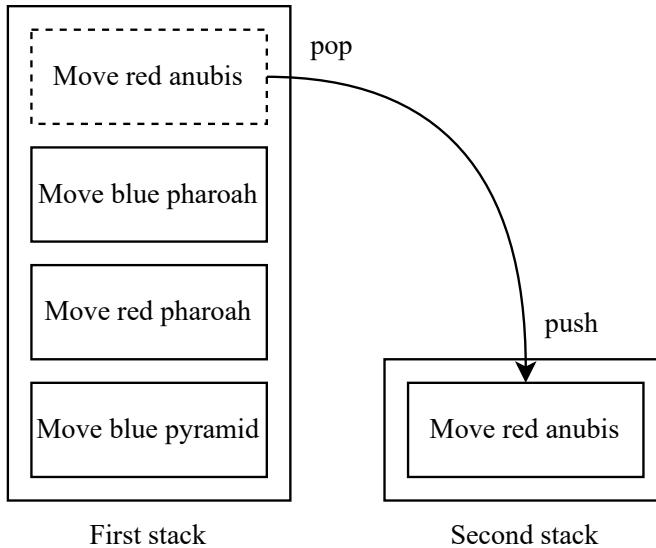


Figure 1.23: *Move red anubis* is undone and pushed onto the second stack

1.4 Classes

I will be using an Object-Oriented Programming (OOP) paradigm for my program. OOP reduces repetition of code, as inheritance can be used to abstract repetitive code into a base class, as shown in my widgets implementation. Testing and debugging classes will make my workflow more efficient. This section documents the base classes I am going to implement in my program.

State

Since there will be multiple screens in my program as demonstrated in Figure 1.1, the State base class will be used to handle the logic for each screen. For each screen, void functions will be inherited and overwritten, each containing their own logic for that specific screen. For example, all screens will call the startup function in Table 1.2 to initialise variables needed for that screen. This polymorphism approach allows me to use another Control class to enable easy switching between screens, without worrying about the internal logic of that screen. Virtual methods also allow methods such as `draw` to be abstracted to the State superclass, reducing code in the inherited subclasses, while allowing them to override the methods and add their own logic.

Method Name	Description
<code>startup</code>	Initialise variables and functions when state set as displayed screen
<code>cleanup</code>	Cleanup any variables and functions when state removed from screen
<code>draw</code>	Draw to display
<code>update</code>	Update any variables for every game tick
<code>handle_resize</code>	Scale GUI when window resized
<code>get_event</code>	Receive pygame events as argument and process them

Table 1.2: Methods for State class

Widget

I will be implementing my own widget system for creating the game GUI. This allows me to fully customise all graphical elements on the screen, and also create a resizing system that adheres to Objective ??. The default pygame rescaling options also simply resize elements without accounting for aspect ratios or resolution, and I could not find a library that suits my needs. Having a bespoke GUI implementation also justifies my use of Pygame over other Python frameworks.

I will be utilising the Pygame sprite system for my GUI. All GUI widgets will be subclasses inheriting from the base Widget class, which itself is a subclass of the Pygame sprite class. Since Pygame sprites are drawn via a spritegroup class, I will also have to create a custom subclass inheriting that as well. As with the State class, polymorphism will allow the spritegroup class to render all widgets regardless of their functionality. Each widget will override their base methods, especially the draw (set_image) method, for their own needs. Additionally, I will use getter and setter methods, used with the `@property` decorator in python, to compute attributes mainly used for resizing widgets. This allows me to expose common variables, and to reduce code repetition.

Method Name	Description
set_image	Render widget to internal image attribute for pygame sprite class
set_geometry	Set position and size of image
set_screen_size	Set screen size for resizing purposes
get_event	Receives pygame events and processes them
screen_size*	Returns screen size in pixels
position*	Returns topleft of widget rect
size*	Returns size of widget in pixels
margin*	Returns distance between border and actual widget image
border_width*	Returns border width
border_radius*	Returns border radius for rounded corners
font_size*	Returns font size for text-based widgets

* represents getter method / property

Table 1.3: Methods for Widget class

I will also employ multiple inheritance to combine different base class functionalities together. For example, I will create a pressable base class, designed to be subclassed along with the widget class. This will provide attributes and methods for widgets that support clicking and dragging. Following Python's Method Resolution Order (MRO), additional base classes should be referenced first, having priority over the base Widget class.

Method Name	Description
get_event	Receives Pygame events and sets current state accordingly
set_state	Sets current Pressable state, called by <code>get_event</code>
set_colours	Set fill colour according to widget Pressable state

Method Name	Description
current_state*	Returns current Pressable state (e.g. hovered, pressed etc.)
* represents getter method / property	

Table 1.4: Methods for example Pressable class

Game

For my game screen, I will be utilising the Model-View-Controller architectural pattern (MVC). MVC defines three interconnected parts, the model processing information, the view showing the information, and the controlling receiving user inputs and connecting the two. This will allow me to decompose the development process into individual parts for the game logic, graphics and user input, speeding up the development process and making testing easier. It also allows me to implement multiple views, for the pause and win screens as well. For MVC, I will have to implement a game model class, a game controller class, and three classes for each view (game, pause, win). Using aggregation, these will be initially connected and handled by the game state class. For the following methods, I have only showed those pertinent to the MVC pattern:

Method Name	Description
get_event	Receives Pygame events and passes them onto the correct part's event handler
handle_game_event	Receives events and notifies the game model and game view
handle_pause_event	Receives events and notifies the pause view
handle_win_event	Receives events and notifies the win view
...	...

Table 1.5: Methods for Controller class

Method Name	Description
process_model_event	Receives events from the model and calls the relevant method to display that information
convert_mouse_pos	Sends controller class information of widget under mouse
draw	Draw information to display
handle_resize	Scale GUI when window resized
...	...

Table 1.6: Methods for View class

Method Name	Description
register_listener	Subscribes method on view instance to an event type, so that the method receives and processes that event everytime <code>alert_listener</code> is called
alert_listener	Sends event to all subscribed instances
toggle_win	Sends event for win view
toggle_pause	Sends event for pause view

Method Name	Description
...	...

Table 1.7: Methods for Model class

Shaders

To use ModernGL with Pygame, I have created classes for each fragment shader, controlled by a main shader manager class. The fragment shader classes will rely on composition: The shader manager creates the fragment shader class; Every fragment shader class takes their shader manager parent instance as an argument, and runs methods on it to produce the final output.

1.4.1 Class Diagram

